Lab 6: Babble – A multi-threaded server (for social networking)

Master M1 MOSIG – Université Grenoble Alpes (Ensimag & UFR IM2AG)

2024

In this lab, you are going to build a *real* multi-threaded application out of a sequential program. This program is rather complex, as it implements a client-server communication protocol. Hence, this lab is not only a chance to practice concurrent programming but also to get used to work with code of significant size.

During the different stages of the project, you will have to identify the concurrency issues related to the use of multiple threads, and to propose appropriate solutions to handle these issues.

1 Important information

- This assignment will be graded.

- The assignment is to be done by groups of at most **2** students.

- Deadline: Friday, November 22, 2024 at 23:59.

- The assignment is to be turned in on Moodle (report + source code).

1.1 Collaboration and plagiarism

You are encouraged to discuss ideas and problems related to this project with the other students. You can also look for additional resources on the Internet. However, we consider plagiarism very seriously. Hence, if any part of your final submission reflects influences from external sources, you must cite these external sources in your report. Also, any part of your design, your implementation, and your report should come from you and not from other students. We will run tests to detect similarities between source codes. Additionally, we will allow ourselves to question you about your submission if part of it looks suspicious, which means that you should be able to explain each line you submitted. In case of plagiarism, your submission will not be graded and appropriate actions will be taken.

1.2 Evaluation of your work

The main points that will be evaluated for this work are (ordered by decreasing priority):

- The understanding of the problem.

- The design of the proposed solution.

- The correctness of the code.

- The quality of the code.

- The number of implemented features.

1.3 Your submission

Your submission is to be turned in on Moodle as an archive named with the last name of the two students involved in the project: `Name1_Name2_lab6.tar.gz`.

The archive will include:

- A **concise** report, either in `txt`, `md` (MarkDown) or `pdf` format[1], which should include the following sections:

  - The name of the participants.
  - For each stage of the project:
    * A short description of each concurrency issue that you have identified.
    * A few words about the solution that you designed to solve it (Note that a simple figure is sometimes better than a long paragraph).
    * Any additional information that you consider necessary to understand your code.
    * A list of tests that you pass successfully, and the known bugs if any.
    * A brief description of the new tests that you have designed, if any (bonus).
  - A feedback section including any suggestions you would have to improve this lab.

- **A version of your source code for <u>each stage</u> of the project** (except for stage 0), in a directory `stage_X` for stage X. **Each version should be self-contained**.

1.4 Expected achievements

Considering the time that is given to you to work on the assignment, here is a scale of our expectations:

- An **acceptable work** is one in which at least stage 1 has been implemented and works correctly.

- A **good work** is one in which:

  - the conceptual questions of stage 2 (steps 2a and 2b) have been correctly answered.
  - the first part of the implementation of stage 2 (step 2c) has been implemented, and works correctly.

- An **excellent work** is one in which, all stages have been implemented and extensively tested.

---

[1]Other formats will be rejected.

1.5   Tentative schedule

For indicative purposes, we provide below a indicative schedule that you should try to follow in order to properly manage the work and time needed for completing the project.

- **By the end of the first week** (i.e., <u>before</u> the second lab session), you should <u>at least</u> have completed Stage 1.

- **By the end of the second week** (i.e., <u>before</u> the third lab session), you should have completed steps 2.a, 2.b, and 2.c.

1.6   About multi-threaded programming

In this lab, you are going to program with `pthreads`. To solve concurrency issues, you are allowed to use any synchronization mechanisms introduced during the lectures. This includes: pthread mutexes, condition variables, semaphores, barriers (see `man pthread_barrier_init`), read-write locks (see `man pthread_rwlock_init`) and atomic operations[2].

2   Overview

This lab is about the implementation of a client-server application providing a social networking service. This service is called Babble. It allows users to register to the service and to publish messages. A user can also follow another user. The publications of a user are only visible to her followers.

2.1   About client-server communication

The goal of this lab is not to have you dealing with issues related to network communication. Hence, the code to communicate between Babble clients and the server is provided to you. In the following, we shortly describe network communication in C, for you to better understand the code you are going to manipulate. Many resources are available on the Internet if you would like more information on the topic[3].

The principal OS abstraction for network communication is the *socket*. It is the endpoint at which a process attaches itself to the network. In this lab, we are working with TCP sockets, which provide connection-oriented communication. It means that a communication session is established between the client and the server before any application data is transferred. Here are the steps used to establish a connection between a server and a client (for more details, consult the man pages of the system calls):

1. A socket is identified by a port number. The port number is a value between 0 and 65535[4].

---

[2]A list of built-in atomic operations for gcc is available at `https://gcc.gnu.org/onlinedocs/gcc-4.4.0/gcc/Atomic-Builtins.html`

[3]See for example: `http://pages.cs.wisc.edu/~suman/courses/640/lectures/sockets.pdf`

[4]Values from 0 to 1023 correspond to *system ports* and require administrator privileges.

Babble uses port number 5656 by default.

Steps run on the server side:

2. A call to `socket()` is used to create a new socket on the server.

3. A call to `bind()` attaches the socket to the selected port.

4. Calling `listen()` informs that we are going to listen for new connections on this socket.

5. Calling `accept()` blocks until a new connection arrives on the socket. Then, it creates a new socket and return a file descriptor for that socket. This file descriptor will be used to read/write data to the connected client.

Steps run on the client side:

6. On the client side, a socket should also first be created using `socket()`.

7. `connect()` is called to connect to the server socket defined by an IP address (or hostname) and the port number the server is listening on. If the client and the server are on the same host, `localhost` can be used as hostname (or IP address `127.0.0.1`). `connect()` also returns a file descriptor.

To communicate over the established connection, `read()` and `write()` operations can be run on the obtained file descriptors. The Babble connections are full-duplex: communication can happen in both directions at the same time.

2.2 The Babble messages

A set of messages are defined for the Babble communication protocol:

- `LOGIN id`: First message to send to register to the service. The `id` has a maximum size of 16 and can only include alphanumeric characters (no white spaces).

- `PUBLISH msg`: To publish a new message of max size 64. The message can include any (ASCII) characters but no white spaces.

- `FOLLOW id`: To follow another registered user. Note that at the time it registers (i.e., upon login), a user only follows herself.

- `TIMELINE`: To get a list of the messages published by the people the user follows since her last `TIMELINE` message. Note than when a user starts following another user, she does not get access to the previously published messages. Note also that if the timeline is big, an information about the total number of messages is sent, but the content of only the last 4 messages is included in the answer.

- `FOLLOW_COUNT`: To get the number of followers of the user.

- `RDV`: Used to check if all the messages[5] sent previously have been processed by the server.

Important: Note that, in the current version of Babble, if a client disconnects and connects again later on with the same `id`, it will be considered as a new user. Also, the `FOLLOW` command only allows following users that are connected at the time the command is run.

2.3 Provided material

You are provided with a sequential implementation of Babble. It includes a Makefile to compile all the source files, and create the executable files for the server and the client (as well as some tests). Feel free to create additional source files to structure your code, but do not forget to update the Makefile accordingly.

The list of provided source file is given below; the names shown in bold font correspond to the ones that you will probably need to modify in order to complete the lab.

- **babble_config.h**: Defines configuration parameters
- **babble_types.h**: Defines some of the main data structures of Babble
- **babble_timeline***: Operations on a timeline
- **babble_registration***: Management of data associated with connected clients
- `babble_communication*`: Utility functions to communicate between the server and the clients
- `babble_utils*`: Other utility functions
- `babble_server.h`: Defines the main functions used by the server
- **babble_server_implem.c**: Implementation of the functions defined in `babble_server.h`
- `babble_server_answers*`: Utility functions for sending answers from the server to the clients.
- `babble_client.h`: Defines the main functions used by the client
- `babble_client_implem.c`: Implementation of the functions defined in `babble_client.h`
- **babble_server.c**: Implementation of the server
- `babble_client.c`: Implementation of the client
- `fastrand.*`: Thread-safe fast random number generator
- `*_test.c`: Code of the provided tests

Important:

- You are free to add new files whenever required for your solution; for example, regarding the commande buffer to be introduced in stage 1 (see §4).

- The provided code (as mentioned in the inlined comments) has some memory leaks. This is a deliberate choice in order to simplify the code base. You are not expected to fix these memory leaks in this lab assignment (but should avoid introducing other ones).

---

[5]More precisely: all messages of any kind (`PUBLISH`, `FOLLOW` ...).

- The provided code includes programs to help you testing your solutions. These aspects are described in Section 9.

- Please do not hesitate to ask for explanations about the provided code by questioning the teaching staff during the lab sessions or through the forum on Moodle.

3 Stage 0: First contact with Babble

3.1 Running the code

It is now time to start working with Babble. As a very first step, we suggest you to try running the server and a client. After compiling the code, the following command can be used to run the server:

```
$ ./babble_server.run
```

To know the options that can be passed to the server executable, run with "-h".

In another terminal, launch the client:

```
$ ./babble_client.run -i my_id
```

Option "-i" is used to specify the identifier of the new client. Use "-h" to learn about other options.

Each client is identified on the server by a unique key which corresponds to a hash of its id. The client receives this key as an acknowledgment of successful login.

The babble client is a primitive console where you can type commands to be sent to the server. For instance, to publish the message "voila", simply type:

```
PUBLISH voila
```

Note that for the sake of simplicity, each command is also represented by a number (1 for PUBLISH) that can be used in the console. Hence, one can also enter:

```
1 voila
```

To find out the number that corresponds to each command, check babble_types.h.

Finally, to terminate the client console, one can use ⒸⓉⓇⓁ + Ⓓ on an empty line to generate an *end-of-file* condition.

You can now play with Babble. However, you will soon notice that this sequential version is rather boring, as it accepts a single client connection at a time. But before starting implementing a multi-threaded version of Babble, let's try to better understand the code.

3.2 Digging into the code

To help you start getting familiar with the code of Babble, we suggest you to answer the following questions:

1. Which part of the code opens the socket where the server is going to listen for new connections?

2. Which part of the code manages new connections on the server side?

3. What are the major steps that are run on the server when it receives a message from the client?

4. Why are `LOGIN` messages managed differently from other commands?

5. What is the purpose of the `registration_table`?

6. How are keys used on the server?

7. How are the answers to client requests implemented on the server?

8. What happens on the execution of a `FOLLOW` request?

9. What happens on the execution of a `PUBLISH` request?

10. What happens on the execution of a `TIMELINE` request?

11. What are the high-level functions provided to communicate over the network? What is their return value?

4 Stage 1: Accepting multiple client connections

In this first stage, we would like to allow multiple clients to be registered to the service at the same time. To do so, we propose to create multiple threads on the server. Namely, we identify two new kinds of threads:

Communication threads. A communication thread is responsible for receiving from the network the messages sent by a client, and to parse these messages to generate commands to be executed.

Executor threads. An executor thread is responsible for executing pending commands and sending answers to clients.

In stage 1, a new communication thread will be created each time a new client connects to the server. On the other hand, a single executor thread will be created and will execute the commands of all clients.

To pass commands created by the communication threads to the executor thread, you should create a buffer (the `command_buffer`), and you should solve the producer-consumer synchronization issues on this buffer.

Figures 1, 2, and 3 provide a basic description of the sequence of tasks to be run by the threads in this stage. Note that all the code required to run these tasks is already provided to you: in the provided version, everything is executed by the main thread.

```
1   Parse input arguments
2   Initialize server data structures (call to server_data_init())
3   Start the executor thread
4   Initialize and open the server socket (call to server_connection_init())
5   while(true){
6      Accept a new connection and get a file descriptor on the newly created
           socket (call to server_connection_accept())
7      Start a new communication thread to handle the new client with the file
           descriptor as argument
8   }
```

Figure 1: Executed by the main server thread in Stage 1

```
1   Recv login message (call to network_recv())
2   Parse the message to get the command (call to parse_command())
3   Process the login command (call to process_command())
4   Send answer to the client (call to send_answer_to_client())
5   while(there are client messages to process){
6      Parse the message to get the command
7      Put the command in the command_buffer
8   }
9   Unregister the client
```

Figure 2: Executed by a communication thread in Stage 1

```
1   while(true){
2      Take a command from the command_buffer
3      Process the command and generate an answer
4      Send answer to the client
5   }
```

Figure 3: Executed by the executor thread in Stage 1

In addition to the producer-consumer problem previously mentioned, another issue is induced by these changes. Namely, there is one data structure that may be accessed concurrently by multiple threads. This is the registration_table. This table is read during the execution of multiple functions involved in the processing of commands and it is updated during the processing of the login command as well as during the unregistration of the client.

This creates a reader-writer problem that you should solve. We strongly recommend you to insert the synchronization code directly in the functions that implement access to the registration_table, that is in file babble_registration.c.

Important: The provided code includes a call to `fastRandomSetSeed(time(NULL) + pthread_self()*100)` in the main server thread. A per-thread pseudo-random number generator is used for testing purposes (see Section 9 for more details). To be able to correctly test the code, this call should be executed during the initialization step of each executor thread. Hence, considering Figure 3 that describes the code of an executor thread, **it means that you should insert this call before line 1**.

5 Step 2: Running multiple executor threads

5.1 Evolution of the server

It seems obvious that the server could be more efficient if there were more than one executor thread.

The goal of this stage is to modify the server design so that N executor threads are run[6].

5.1.1 Step 2.a: Identification of concurrency issues

The goal of this step is to identify the potential concurrency issues that are introduced by the use of multiple executor threads. To simplify the problem, we will not consider the `RDV command`.

We advise you to start by describing precisely, for each type of request, the sequence of actions that must be performed to execute a request. More precisely, you should consider the following request types: `FOLLOW`, `PUBLISH`, `TIMELINE`.

Based on this, you should be able to identify the data structures and actions that will require specific care when multiple executor threads are running concurrently. *Describe the issues that you have identified in your report.*

Note: For this step, you are <u>not</u> asked to design or implement solutions.

5.1.2 Step 2.b: Impact of the `RDV` command

Revisit the previous question now taking into account the existence of the `RDV` command. More precisely, handling this command type introduces at least one new problem, in addition to the concurrency issues identified in Step 2.a. *Describe the new problem(s) in your report*.

5.1.3 Step 2.c: Implementation Part 1

Based on the problems identified in steps 2a and 2b, you are asked to find the appropriate solutions and to implement the new version of the server where multiple executor threads are used.

For simplification, in this step only consider the problems identified in Step 2.a[7].

---

[6]`BABBLE_EXECUTOR_THREADS` constant defined in `babble_config.h`.
[7]Since the problem of `RDV` is not properly handled yet, tests with streaming will not work in this step

5.1.4 Step 2.d: Implementation Part 2

Address the requirements of the RDV command in your implementation (Identified in step 2.b).

5.2 Testing the new version

Test your new server with the provided test in normal execution mode. In a second step, test with streaming: having a correct solution for this case might be more subtle.[8]

6 Stage 3: Performance improvements

In this stage, we consider design changes allowing to improve the performance of the Babble service. We will go through two design iterations, the second one being more efficient but also more complex.

Notes:

- You are not required to demonstrate empirically (i.e., with experiments/measurements) that these new designs do indeed improve performance. This would require additional efforts (for example, to make sure that the studied machine/system is properly configured to achieve such improvements) – very interesting questions but out of the scope of this lab work.

- You are simply asked to provide a correct and meaningful implementation (by "meaningful", we mean a that this implementation should be able to bring performance gains on some situations and some appropriate machines/systems).

- Nonetheless, we encourage you to try running a few experiments to observe the performance improvements achieved by your optimized versions. Beyond your personal laptop, you can for example try running performance comparisons on one of the servers of the University (see Section 9 for more details).

The performance optimization to be considered in this stage is related to the interactions between the communication threads and the executor threads. More precisely, the command buffer may become a bottleneck if many threads are concurrently trying to produce/retrieve new requests in it. In other words, we may reach a situation in which most threads are blocked, waiting to obtain exclusive access to the buffer. A way to mitigate this issue is to use several instances of the command buffer instead of a single one. We will assume that the number of buffer instances is defined statically (see the constant BABBLE_PRODCONS_NB in babble_config.h.

In addition to your code, your are asked to summarize, in your report, the main principles of the approach that you propose for each step.

---

[8]Please see §9 for details about tests in normal mode and streaming mode.

6.1 Step 3.a: Static version

You are asked to imagine and implement a strategy for the interactions between the (communication and execution) threads and the buffers instances.

In this version, the command dispatching strategy between the different buffer instances is defined statically. By this, we mean that the algorithms used by the various threads to choose which buffer to use are simple and do not take into account the current usage of each buffer (number of empty slots in the buffer, …).

6.2 Step 3.b: Advanced version

In practice, for various reasons, some requests may take more time to be executed than others. As a consequence, a static request dispatching strategy between the buffer instances (like in step 3.a) may result in load imbalances, i.e., with some threads remaining mostly idle and/or some buffers becoming a bottleneck.

The goal of this new version is to take into account this problem by designing a more advanced approach. For that matter, you can consider one or several of the following ideas:

- Using "batching": that is, allowing an executor thread to retrieve multiple commands from a buffer at once.

- Designing a request dispatching approach that can adjust to the processing speed of the different threads and the remaining capacity of the different buffers.

- Any other idea that you may have (but please contact the teaching staff ahead of time to validate the feasibility beforehand).

7 Stage 4 (Bonus): Dealing with celebrities

The standard way of handling PUBLISH requests in Babble is the following: upon the processing of a PUBLISH command, the new message is written in the timeline of all the followers of the publisher.

This creates a scalability issue for users that have a huge number of followers (celebrities). For celebrities, it is more efficient to maintain a list of publications (one list per celebrity) where all the publications of the celebrity are stored. Then, on a TIMELINE command, a client should not only check its local timeline but it should also collect new publications in the lists of publications of the celebrities it follows.

Implement a new version of the Babble server where clients are handled as *celebrities* when their number of followers surpasses a threshold.

8 Some advice on coding patterns

> To avoid some common mistakes, **we strongly recommend to follow the two guidelines below when crafting your code**:
>
> - **Be mindful of always releasing the lock(s) taken by a thread**, even when you must leave a function early (due to a code shortcut dealing with an edge case or an error situation).
>
> - To avoid concurrency bugs and make your code more readable and maintainable, it is better to **put the synchronization code that protects a given data structure inside the wrapper functions used to manipulate this structure** rather than spread it in various places throughout the code (as already suggested in the case of the registration table for Stage 1).

9 Testing the code

9.1 General remarks about tests

Testing multi-threaded applications is a very complex task. As multi-threading implies non-determinism, successfully running a test does not mean your program is correct. On the other hand, failing to run a test means that your program is not correct (assuming that the test itself is correct).

To have more chances to detect bugs, it can be good to test your code under *extreme configurations*. An extreme configuration can be, for instance, when the size of the buffer used for the producer-consumer problem is very small.

Another technique to increase the chances of detecting bugs to introduce random artificial delays in some part of the code, in order to influence the scheduler and produce more diverse thread interleavings (see §9.3).

9.2 Overview

It is hard to put the server under stress condition by simply using the client console. This is why you are provided with three additional programs on client side to be used for testing :

follow_test: This test creates two threads (corresponding to two clients called `PUB` and `TIM`). `TIM` follows `PUB`. `PUB` continuously publishes new messages while `TIM` sends `TIMELINE` requests to get the new publications. At the end of the test, the number of publications received by `TIM` should correspond to the number of publications of `PUB`.

stress_test: This test creates several threads, each corresponding to a client. First, each client follows all other clients. At the end of the step, the test checks that the answers to `FOLLOW_COUNT` are consistent. Then, each thread publishes a sequence of messages. At the end of the second step, the test verifies that the answers to `TIMELINE` are consistent. After-wards, half of the initial clients are disconnected and are replaced with new clients, and a new testing phase (similar to the first one) is launched.

performance_test: This test measures the throughput of the server in terms of number of messages processed per second. It creates several threads, each corresponding to a client. Each client tries publishing as many messages as possible during a predefined amount of time.

For each test, use "-h" to get information about the options. Do not hesitate to run the tests with different input parameters.

9.3 Random delays

The provided code also includes another feature aimed at facilitating randomized testing. When enabled this features inserts random delays (between 0 and 10 milliseconds in the processing of some types of commands (PUBLISH, FOLLOW, TIMELINE). This is useful to simulate more realistic scenarios and, mainly, to increase the probability to observe more "subtle" thread interleavings and hence to catch more bugs related to the design and implementation for the concurrency support within Babble. To enable this feature, you must launch the server with the following option: -r.

Warning: For random delays to work properly, one should pay attention to the initialization of the random number generator (see Section 4).

Note that, depending on the stage and scenario that you consider, it is sometimes more appropriate (in order to exhibit bugs more easily) to enable the random delays or, conversely, to disable them (see Section 9.5)

9.4 Streaming

The three provided tests have an option ("-s") to activate streaming. In streaming mode, the client does not wait for an answer for the commands that do not explicitly need it, that is, FOLLOW and PUBLISH. This allows stressing even more the server by having clients continuously sending new requests (to avoid totally overloading the server, in the provided test, there is a delay of 0.1 ms between two requests – except for the performance test). Do not hesitate to test also with streaming activated[9].

Note that the use of streaming required introducing the RDV command for testing purposes. The RDV command asks for an acknowledgment by the server that is sent only when all the preceding requests of the client have been processed. Without this, it would be impossible for a client to know if all its PUBLISH requests have been handled in streaming mode.

Note also that the streaming mode is implemented by prefixing the request sent to the client with "S". As such, you can also test streaming in the console with:

```
S PUBLISH voila
```

Warning: Tests in streaming mode are only relevant/useful/valid for Stage 2 (step 2.d) and onwards.

---

[9]Running the performance_test with streaming and random delays before stage 2 is not recommended as it might take a lot of time for the server to process the large amount of messages sent.

9.5   About performance tests

Please find below some recommendations about running the `performance_test` program:

- Except for Stage 3, it is more meaningful to run the tests while activating random delays on the server side.

- To have a good performance evaluation for Stage 3, it is better to avoid displaying messages in the console for each processed message on server side. It is also better to activate streaming.

- In general, it is not very useful to run a performance test for a long period of time (a few seconds is enough).

- From Stage 2, to determine the maximum achievable throughput for the server, it is important to run tests with increasing number of clients until reaching the saturation point.